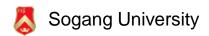
Computer Architecture

Chap 1. Computer Abstractions and Technology

Hyuk-Jun Lee, PhD

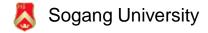
Dept. of Computer Science and Engineering Sogang University Seoul, Korea

Email: hyukjunl@sogang.ac.kr

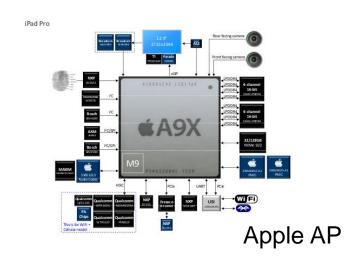


The Computer Revolution

- Progress in computer technology
 - Underpinned by Moore's Law
- Makes novel applications feasible
 - Mobile Computing: Cell phones
 - Bioengineering: Human genome project
 - Data Mining: World Wide Web, Search Engines
 - Machine Learning: Deep learning
- Computers are pervasive

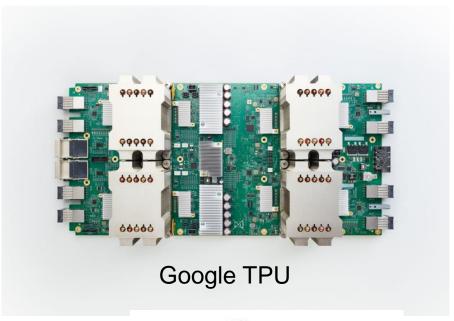


CPUs

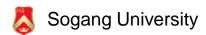






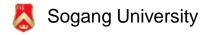




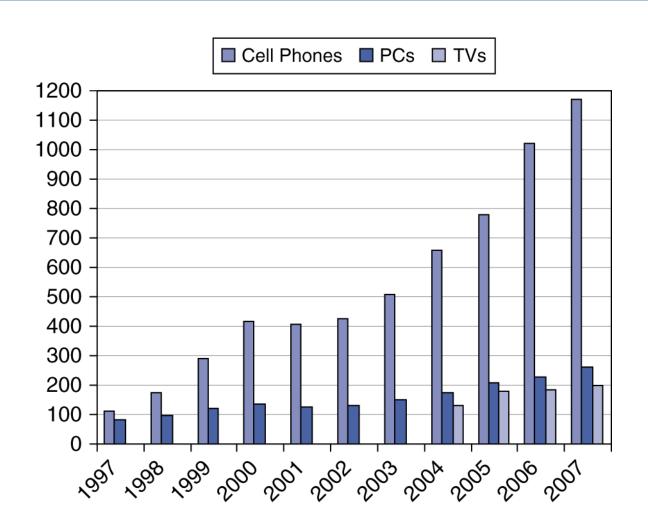


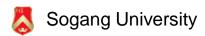
Classes of Computers

- Desktop computers
 - General purpose, variety of software
 - Subject to cost/performance tradeoff
- Server computers
 - Network based
 - High capacity, performance, reliability
 - Range from small servers to building sized
- Embedded computers
 - Hidden as components of systems
 - Stringent power/performance/cost constraints



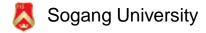
The Processor Market





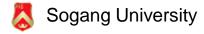
What You Will Learn

- How programs are translated into the machine language
 - And how the hardware executes them
- The hardware/software interface
- What determines program performance
 - And how it can be improved
- How hardware designers improve perfor mance
- What is parallel processing



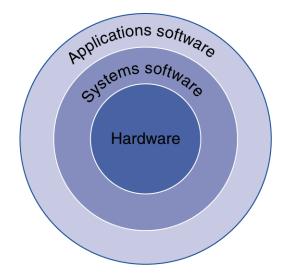
Understanding Performance

- Algorithm
 - Determines number of operations executed
- Programming language, compiler, architecture
 - Determine number of machine instructions executed per operation
- Processor and memory system
 - Determine how fast instructions are executed
- I/O system (including OS)
 - Determines how fast I/O operations are executed



Below Your Program

- Application software
 - Written in high-level language
- System software
 - Compiler: translates HLL code to machine code
 - Operating System: service code
 - Handling input/output
 - Managing memory and storage
 - Scheduling tasks & sharing resources
- Hardware
 - Processor, memory, I/O controllers

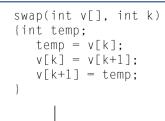


Levels of Program Code

- High-level language
 - Level of abstraction closer to problem domain
 - Provides for productivity and portability
- Assembly language
 - Textual representation of instructions
- Hardware representation
 - Binary digits (bits)
 - Encoded instructions and data

High-level language program (in C)

Assembly language program (for MIPS)

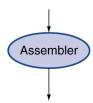


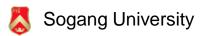
Compiler

swap:

muli \$2, \$5,4
add \$2, \$4,\$2
lw \$15, 0(\$2)
lw \$16, 4(\$2)
sw \$16, 0(\$2)
sw \$15, 4(\$2)


\$31




Binary machine language program (for MIPS) 

Components of a Computer

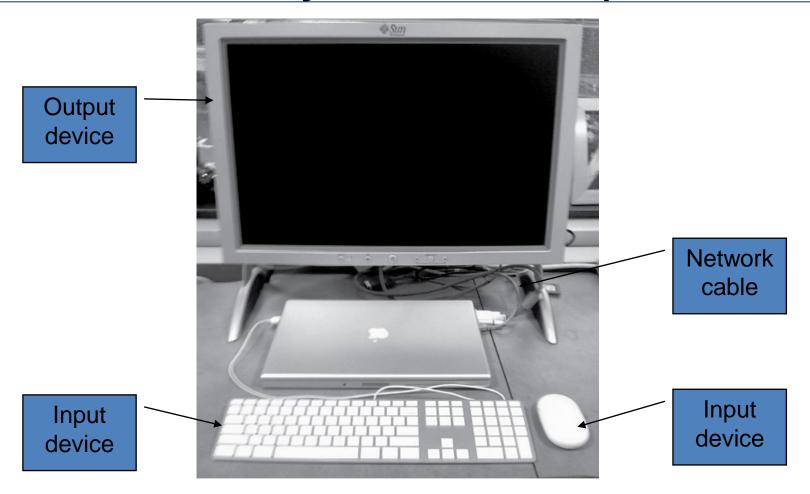
The BIG Picture

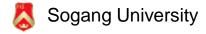


- Same components for all kinds of computer
 - Desktop, server, embedded
- Input/output includes
 - User-interface devices
 - Display, keyboard, mouse
 - Storage devices
 - Hard disk, CD/DVD, flash
 - Network adapters
 - For communicating with other computers



Anatomy of a Computer

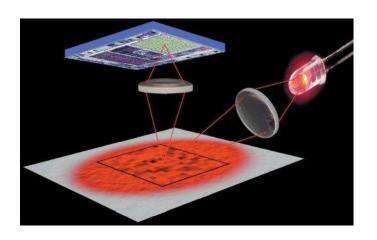




Anatomy of a Mouse


- Optical mouse
 - LED illuminates desktop
 - Small low-res camera
 - Basic image processor
 - Looks for x, y movement
 - Buttons & wheel
- Supersedes roller-ball mechanical mouse







Through the Looking Glass

- LCD screen: picture elements (pixels)
 - Mirrors content of frame buffer memory

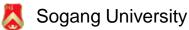




Opening the Box

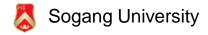







Inside the Processor (CPU)


- Datapath: performs operations on data
- Control: sequences datapath, memory, ...
- Cache memory
 - Small fast SRAM memory for immediate access to data



Inside the Processor

AMD Barcelona: 4 processor cores

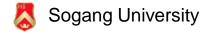




Abstractions

The BIG Picture

- Abstraction helps us deal with complexity
 - Hide lower-level detail
- Instruction set architecture (ISA)
 - The hardware/software interface
- Application binary interface
 - The ISA plus system software interface
- Implementation
 - The details underlying and interface



A Safe Place for Data


- Volatile main memory
 - Loses instructions and data when power off
- Non-volatile secondary memory
 - Magnetic disk
 - Flash memory
 - Optical disk (CDROM, DVD)











Networks


- Communication and resource sharing
- Local area network (LAN): Ethernet
 - Within a building
- Wide area network (WAN: the Internet)
- Wireless network: WiFi, Bluetooth





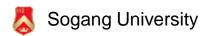
Technology Trends

- Electronics technology continues to evolve
 - Increased capacity and performance
 - Reduced cost




DRAM capacity


Year	Technology	Relative performance/cost
1951	Vacuum tube	1
1965	Transistor	35
1975	Integrated circuit (IC)	900
1995	Very large scale IC (VLSI)	2,400,000
2005	Ultra large scale IC	6,200,000,000





Defining Performance

Which airplane has the best performance?










Response Time and Throughput

- Response time
 - How long it takes to do a task
- Throughput
 - Total work done per unit time
 - e.g., tasks/transactions/... per hour
- How are response time and throughput affected by
 - Replacing the processor with a faster version?
 - Adding more processors?
- We'll focus on response time for now...

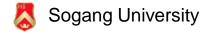


Relative Performance

- Define Performance = 1/Execution Time
- "X is n time faster than Y"

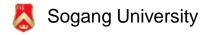
Performanc e_x/Performanc e_y

- = Execution time $_{Y}$ /Execution time $_{X}$ = n
- Example: time taken to run a program
 - 10s on A, 15s on B
 - Execution Time_B / Execution Time_A
 = 15s / 10s = 1.5
 - So A is 1.5 times faster than B




Measuring Execution Time

- Elapsed time
 - Total response time, including all aspects
 - Processing, I/O, OS overhead, idle time
 - Determines system performance
- CPU time
 - Time spent processing a given job
 - Discounts I/O time, other jobs' shares
 - Comprises user CPU time and system CPU time
 - Different programs are affected differently by CPU and system performance

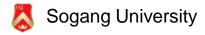


CPU Clocking

Operation of digital hardware governed by a constant-rate clock



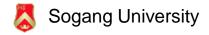
- Clock period: duration of a clock cycle
 - e.g., $250ps = 0.25ns = 250 \times 10^{-12}s$
- Clock frequency (rate): cycles per second
 - e.g., $4.0GHz = 4000MHz = 4.0 \times 10^9Hz$



CPU Time

```
CPU Time = CPU Clock Cycles \times Clock Cycle Time = \frac{\text{CPU Clock Cycles}}{\text{Clock Rate}}
```

- Performance improved by
 - Reducing number of clock cycles
 - Increasing clock rate
 - Hardware designer must often trade off clock rate against cycle count



CPU Time Example

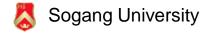
- Computer A: 2GHz clock, 10s CPU time
- Designing Computer B
 - Aim for 6s CPU time
 - Can do faster clock, but causes $1.2 \times \text{clock}$ cycles
- How fast must Computer B clock be?

Clock Rate_B =
$$\frac{\text{Clock Cycles}_{\text{B}}}{\text{CPU Time}_{\text{B}}} = \frac{1.2 \times \text{Clock Cycles}_{\text{A}}}{6\text{s}}$$

Clock Cycles_A = CPU Time_A × Clock Rate_A

$$= 10\text{s} \times 2\text{GHz} = 20 \times 10^{9}$$

Clock Rate_B = $\frac{1.2 \times 20 \times 10^{9}}{6\text{s}} = \frac{24 \times 10^{9}}{6\text{s}} = 4\text{GHz}$



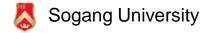
Instruction Count and CPI

Clock Cycles = Instruction Count × Cycles per Instruction

CPU Time = Instruction Count × CPI × Clock Cycle Time


$$= \frac{Instruction Count \times CPI}{Clock Rate}$$

- Instruction Count for a program
 - Determined by program, ISA and compiler
- Average cycles per instruction
 - Determined by CPU hardware
 - If different instructions have different CPI
 - Average CPI affected by instruction mix



CPI Example

- Computer A: Cycle Time = 250ps, CPI = 2.0
- Computer B: Cycle Time = 500ps, CPI = 1.2
- Same ISA
- Which is faster, and by how much?



CPI in More Detail

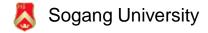
 If different instruction classes take different numbers of cycles

$$Clock \ Cycles = \sum_{i=1}^{n} (CPI_{i} \times Instructio \ n \ Count_{i})$$

Weighted average CPI

$$CPI = \frac{Clock \ Cycles}{Instructio \ n \ Count} = \sum_{i=1}^{n} \left(CPI_{i} \times \frac{Instructio \ n \ Count}{Instructio \ n \ Count} \right)$$

Relative frequency



CPI Example

 Alternative compiled code sequences using instructions in classes A, B, C

Class	А	В	С	
CPI for class	1	2	3	
IC in sequence 1	2	1	2	
IC in sequence 2	4	1	1	

- Sequence 1: IC = 5
 - Clock Cycles= 2×1 + 1×2 + 2×3= 10
 - Avg. CPI = 10/5 = 2.0

- Sequence 2: IC = 6
 - Clock Cycles= 4×1 + 1×2 + 1×3= 9
 - Avg. CPI = 9/6 = 1.5

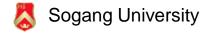


Performance Summary


The BIG Picture

$$CPU \ Time = \frac{Instructio \ ns}{Program} \times \frac{Clock \ cycles}{Instructio \ n} \times \frac{Seconds}{Clock \ cycle}$$

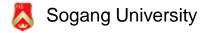
- Performance depends on
 - Algorithm: affects IC, possibly CPI
 - Programming language: affects IC, CPI
 - Compiler: affects IC, CPI
 - Instruction set architecture: affects IC, CPI, T_c



Power Trends



In CMOS IC technology
 Dynamic power = ½ C V² F
 Static power = IV where I is leakage current

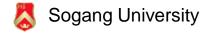


Reducing Dynamic Power


- Suppose a new CPU has
 - -85% of capacitive load of old CPU
 - 15% voltage and 15% frequency reduction

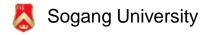
$$\frac{P_{\text{new}}}{P_{\text{old}}} = \frac{C_{\text{old}} \times 0.85 \times (V_{\text{old}} \times 0.85)^2 \times F_{\text{old}} \times 0.85}{C_{\text{old}} \times V_{\text{old}}^2 \times F_{\text{old}}} = 0.85^4 = 0.52$$

- The power wall
 - We can't reduce voltage further
 - We can't remove more heat
- How else can we improve performance?



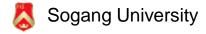
Uniprocessor Performance






Multiprocessors


- Multicore microprocessors
 - More than one processor per chip
- Requires explicitly parallel programming
 - Compare with instruction level parallelism
 - Hardware executes multiple instructions at once
 - Hidden from the programmer
 - Hard to do
 - Programming for performance
 - Load balancing
 - Optimizing communication and synchronization




Manufacturing ICs



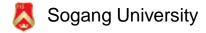
Yield: proportion of working dies per wafer



AMD Opteron X2 Wafer



- X2: 300mm wafer, 117 chips, 90nm technology
- X4: 45nm technology



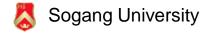
Integrated Circuit Cost

Cost per die =
$$\frac{\text{Cost per wafer}}{\text{Dies per wafer} \times \text{Yield}}$$

Dies per wafer $\approx \text{Wafer area/Die area}$


Yield = $\frac{1}{(1+(\text{Defects per area} \times \text{Die area/2}))^2}$


- Nonlinear relation to area and defect rate
 - Wafer cost and area are fixed
 - Defect rate determined by manufacturing process
 - Die area determined by architecture and circuit design



SPEC CPU Benchmark

- Programs used to measure performance
 - Supposedly typical of actual workload
- Standard Performance Evaluation Corp (SPEC)
 - Develops benchmarks for CPU, I/O, Web, ...
- SPEC CPU2006
 - Elapsed time to execute a selection of programs
 - Negligible I/O, so focuses on CPU performance
 - Normalize relative to reference machine
 - Summarize as geometric mean of performance ratios
 - CINT2006 (integer) and CFP2006 (floating-point)

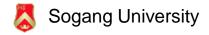




CINT2006 for Opteron X4 2356

Name	Description	IC×10	CPI	Tc (ns)	Exec time	Ref time	SPECratio
perl	Interpreted string processing	2,118	0.75	0.40	637	9,777	15.3
bzip2	Block-sorting compression	2,389	0.85	0.40	817	9,650	11.8
gcc	GNU C Compiler	1,050	1.72	0.47	24	8,050	11.1
mcf	Combinatorial optimization	336	10.00	0.40	1,345	9,120	6.8
go	Go game (AI)	1,658	1.09	0.40	721	10,490	14.6
hmmer	Search gene sequence	2,783	0.80	0.40	890	9,330	10.5
sjeng	Chess game (AI)	2,176	0.96	0.48	37	12,100	14.5
libquantum	Quantum computer simulation	1,623	1.61	0.40	1,047	20,720	19.8
h264avc	Video compression	3,102	0.80	0.40	993	22,130	22.3
omnetpp	Discrete event simulation	587	2.94	0.40	690	6,250	9.1
astar	Games/path finding	1,082	1.79	0.40	773	7,020	9.1
xalancbmk	XML parsing	1,058	2.70	0.40	1,143	6,900	6.0
Geometric mean /						11.7	

High cache miss rates



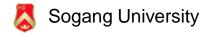
SPEC Power Benchmark

- Power consumption of server at different workload levels
 - Performance: ssj_ops/sec
 - Power: Watts (Joules/sec)

Overall ssj_ops per Watt =
$$\left(\sum_{i=0}^{10} ssj_ops_i\right) / \left(\sum_{i=0}^{10} power_i\right)$$

SPECpower_ssj2008 for X4

Target Load %	Performance (ssj_ops/sec)	Average Power (Watts)
100%	231,867	295
90%	211,282	286
80%	185,803	275
70%	163,427	265
60%	140,160	256
50%	118,324	246
40%	920,35	233
30%	70,500	222
20%	47,126	206
10%	23,066	180
0%	0	141
Overall sum	1,283,590	2,605
∑ssj_ops/ ∑power		493



Pitfall: Amdahl's Law

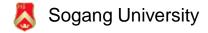
 Improving an aspect of a computer and expecting a proportional improvement in overall performance

$$T_{\text{improved}} = \frac{T_{\text{affected}}}{\text{improvemen t factor}} + T_{\text{unaffected}}$$

- Example: multiply accounts for 80s/100s
 - How much improvement in multiply performance to get 5× overall?

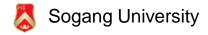
$$20 = \frac{80}{n} + 20$$
 • Can't be done!

Corollary: make the common case fast



Fallacy: Low Power at Idle

- Look back at X4 power benchmark
 - At 100% load: 295W
 - At 50% load: 246W (83%)
 - At 10% load: 180W (61%)
- Google data center
 - Mostly operates at 10% 50% load
 - At 100% load less than 1% of the time
- Consider designing processors to make p ower proportional to load

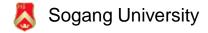


Pitfall: MIPS as a Performance Metric

- MIPS: Millions of Instructions Per Second
 - Doesn't account for
 - Differences in ISAs between computers
 - Differences in complexity between instructions

$$\begin{split} \text{MIPS} &= \frac{\text{Instructio n count}}{\text{Execution time} \times 10^6} \\ &= \frac{\text{Instructio n count}}{\frac{\text{Instructio n count} \times \text{CPI}}{\text{Clock rate}}} = \frac{\text{Clock rate}}{\text{CPI} \times 10^6} \end{split}$$

CPI varies between programs on a given CPU



Concluding Remarks

- Cost/performance is improving
 - Due to underlying technology development
- Hierarchical layers of abstraction
 - In both hardware and software
- Instruction set architecture
 - The hardware/software interface
- Execution time: the best performance measure
- Power is a limiting factor
 - Use parallelism to improve performance

